

The Measure Makes the Finding

How measurement, design, and framing shape what we conclude about race, democracy, and punishment in the U.S. states

Miura Meng

2026-06-17

Abstract

Conclusions about racial inequality often depend more on how we measure and frame a question than on the evidence itself. That measurement and framing are not neutral is by now well established; what is rarer is a compact, concrete demonstration of how far the answer can move, and how easily the moves compound, when a single high-stakes question is carried through the ordinary choices an analyst makes. This paper offers such a demonstration, using race, democracy, and punishment across the U.S. states. It walks the case through four routine decisions. The choice of inequality measure can make more democratic states look either more or less racially unequal, because the popular Black/White ratio moves with its White denominator, which falls steeply in more democratic states while Black imprisonment changes much less. The choice between comparing states and following one state over time turns a striking cross-sectional pattern into a small within-state one. A clustering method returns a tidy five-group typology that looks more solid than the data warrant. And Florida's 2018 Amendment 4 reads as a sweeping expansion of voting rights or as a non-event, depending on whether one reads its text or its effect; a follow-on law (SB7066) requiring payment of court fines left turnout unmoved. In each case a small decision settles the answer. The substantive findings are secondary to that methodological point. Methods include multilevel models, group-based trajectory modeling, and synthetic control.

Introduction

A finding rarely survives retelling intact. As a result moves from the analysis to the abstract to the headline, its conditions fall away, and a claim that began as tentative comes to sound certain. The same number, framed one way, points to individual responsibility; framed another way, it indicts institutions, even though nothing underneath has changed. Quantitative social science is especially exposed to this, because a single summary figure can quietly carry all the choices that produced it.

This paper takes that vulnerability as its subject. On the surface it asks a substantive question: across the U.S. states, does a healthier democracy go together with less racial inequality in punishment, and can restoring the vote to people with felony convictions change anything? Underneath, though, the argument is methodological, because at every step the answer turns on a decision that is easy to make without noticing. Which measure of inequality should be reported? Should states be compared with one another, or each state with itself over time? How many groups should be read into the data? And is a reform “large” by its text or by its effect? Each of these choices, it turns out, can flip the conclusion.

The paper proceeds in four demonstrations, and each one revises the one before it.

- The first measures racial inequality across states. By one common measure, more democratic states look more unequal; by three others, they look the same or better.
- The second compares each state with itself over time, and the cross-state result largely dissolves.
- The third sorts states into trajectory types. The types are real enough to describe, but the method reads in more structure than the data can support.

- The fourth treats Florida’s 2018 voting-rights reform as a natural experiment. By its text it freed 1.4 million voters; by its effect it changed almost nothing.

The point is not any single number, but the pattern across them: measurement, design, and framing decide what we are able to conclude, and careful choices keep overturning the answer that looked settled. The case, race and democracy and punishment, matters in its own right, yet here it mainly serves to make that methodological lesson concrete.

Background

Measurement, framing, and what counts as a finding

A claim about inequality is never simply a reading of the evidence. It is also a product of the choices made in turning evidence into a number, and a number into a sentence. Scholars writing under banners such as data feminism and QuantCrit have pressed this point in recent years, arguing that data are not neutral and that statistics, far from speaking for themselves, carry the assumptions of whoever collected, classified, and summarized them (D’Ignazio & Klein, 2020; Gillborn, Warmington, & Demack, 2018). The complaint is not that quantification is worthless but that it hides its own premises. Measurement of inequality is a clear example: faced with the same data, an analyst can often choose the measure that best supports a preferred conclusion, and the decision between an absolute and a relative measure alone can send a trend in opposite directions (Kjellsson, Gerdtham, & Petrie, 2015; Moonesinghe & Beckles, 2015). Framing does comparable work even when the numbers are held fixed. Describing a gap as one of “achievement” rather than “opportunity,” for instance, shifts where readers locate its causes, away from structures and toward individuals, with-

out altering a single statistic (Quinn, 2025). And as a result travels outward from the model that produced it, the qualifications that gave it meaning tend to drop away, so that a conditional finding begins to circulate as a settled fact (van der Bles, van der Linden, Freeman, & Spiegelhalter, 2020; Merry, 2016).

As a principle, none of this is new. What it has lacked is a compact, concrete demonstration: a single consequential question carried through the ordinary choices an analyst makes, with the answer shown to move at each step. That is what this paper sets out to provide. It takes one case and walks it through four routine decisions, treating each decision as part of the analysis rather than as a preliminary to it. The contribution, then, is not the insight that measurement and framing matter, which is by now well established, but the worked example that shows how much they matter, and how readily they compound, once a real claim about racial inequality is at stake.

Measuring racial disparity in incarceration

The most common summary of racial disparity in imprisonment is a single ratio, the Black imprisonment rate divided by the white rate. Its appeal is obvious, but it is a quotient, and a quotient can change because its numerator changes, because its denominator changes, or because of both at once. A state can post a high ratio because it imprisons Black residents at an unusually high rate, or because it imprisons white residents at an unusually low one, and these are very different social facts that the ratio reports identically. Criminologists have long shown that conclusions about racial disproportionality hinge on such choices: on how much of a gap is attributed to differential involvement rather than to downstream processing (Beck & Blumstein, 2018), on which benchmark population is placed in the denominator (Cesario, Johnson, & Terrill, 2019), and even on the

data source and racial-classification scheme, which can widen or narrow the measured Black/White gap (Sabol, Johnson, & Lynch, 2025). Rather than commit to one measure in advance, this paper carries three in parallel, the ratio, the absolute Black rate, and the gap between the Black and white rates, and treats their disagreement as part of the evidence.

The case: state democracy, disenfranchisement, and Florida

State democracy can be measured and tracked. The State Democracy Index scores each state on features such as electoral participation, the fairness of district maps, and the accessibility of voting, and those scores move over time, including downward in a subset of states (Grumbach, 2022, 2023). Criminal disenfranchisement, meanwhile, is one of the oldest and most racially loaded boundaries on the vote. Many states wrote or expanded these laws in the decades after the Civil War, some in terms understood at the time as a way to limit Black political power, and the legacy persists: roughly four million Americans remain barred from voting by a felony record, with Black citizens barred at several times the rate of others (Manza & Uggen, 2006; The Sentencing Project, 2024). Florida concentrated this history into a single recent episode. In 2018 its voters passed Amendment 4, restoring the vote to about 1.4 million people upon completion of their sentences. Within months, however, the legislature passed SB7066, which redefined “completion” to require paying all outstanding court fines and fees first, a condition critics likened to a poll tax and a federal appeals court allowed to stand in 2020.

Data and measures

The analysis combines three public datasets, each at the state-year level. The State Democracy Index 2.0 (Grumbach & Bitton, 2024) scores the fifty states from 2000 to 2023, with higher values meaning more democratic. The Vera Institute’s Incarceration Trends provides race-specific imprisonment rates per 100,000; the Black and white rates used here run from 1990 to 2022. The UF Election Lab turnout series (McDonald, 2023) gives voting-age-population (VAP) turnout for each even-year election, along with the number of people disenfranchised by a felony record.


Because racial inequality has no single definition, the paper follows three measures side by side: the Black/White rate ratio, the absolute Black rate, and the gap between the Black and white rates. Over this period the national ratio sits near a median of 5.4. For the natural experiment the outcome is VAP turnout, whose denominator is the voting-age population and so does not move mechanically when people are re-enfranchised, and the first-stage quantity is the share of the voting-age population disenfranchised by a felony record.

Demonstration 1: the measure decides the sign

Begin with the question in its most direct form. For each state I average the Black/White imprisonment ratio and the State Democracy Index over 2016 to 2020, and plot one against the other (Figure 1). The line slopes upward: the more democratic a state is, the wider its Black/White ratio tends to be. Read at face value, this is an uncomfortable result, because it seems to say that democratic health and racial fairness in punishment are in tension, and that the states doing democracy best are the ones treating Black and white residents most unequally behind bars.

Weaker state democracy, wider racial incarceration gap?

Each point = one U.S. state, averaged 2016-2020. y = Black-to-White prison-rate ratio.




Sources: Grumbach & Bitton, State Democracy Index 2.0 (UC Berkeley); Vera Institute, Incarceration Trends.
Descriptive association only - not causal.

Figure 1: Across states, averaged over 2016 to 2020, the Black/White incarceration ratio rises with the State Democracy Index.

Before accepting that, it is worth asking whether the result is a fact about the world or a fact about the ratio. The test is simple. I hold the horizontal axis fixed, so the measure of democracy never changes, and I substitute different measures of “racial inequality” on the vertical axis (Figure 2). If the relationship is real, it should survive the substitution. It does not. The Black/White ratio correlates positively with democracy, at about +0.26, which is the pattern we just saw. The absolute Black imprisonment rate, though, correlates slightly negatively, at about -0.14 , so by that measure more democratic states imprison Black residents at marginally lower rates rather than higher ones. The gap between the Black and white rates is essentially flat, at about -0.04 . And the absolute white rate correlates strongly negatively, at about -0.34 . Four measures of the supposedly same inequality, and they do not merely differ in strength; they point in different directions.

One democracy axis, four ways to measure 'racial inequity' - the sign flips

Each point = one U.S. state, averaged 2016-2020. x = State Democracy Index 2.0 (higher = more democratic).



Sources: Grumbach & Bitton SDI 2.0; Vera Institute Incarceration Trends. Descriptive / cross-sectional only.

Figure 2: The same democracy axis against four measures of racial inequity. The sign of the relationship flips depending on the measure.

The reason comes into focus once we look at what the ratio is made of. It is the Black rate divided by the white rate, so its value depends as much on the denominator as on the numerator. More democratic states, it turns out, imprison white residents at strikingly low rates, and that small denominator does the arithmetic work, inflating the ratio even in a state where Black imprisonment is no higher than average. The South illustrates the same mechanism in reverse. Southern states post lower ratios, but not because they treat Black and white residents more evenly; they imprison white residents heavily too, and the large white denominator pulls the ratio down. The regional averages put numbers on this. White imprisonment runs about 230 per 100,000 in the more democratic Northeast and about 535 in the South, more than twice as high, while the Black rate stays

within a much narrower band, roughly 1,460 to 2,060 across the four regions. The ratio therefore runs highest on average across the low-white-rate Northeast and lowest across the high-white-rate South, even though the South is the less democratic region. The measure reported most often, then, is also the one most exposed to movement that has nothing to do with how Black people are treated, since a state can climb the ratio simply by incarcerating fewer white people. Judged by the absolute Black rate, more democratic states look slightly better; judged by the gap, they look no different; judged by the ratio, they look worse. Nothing about the data has changed between those three sentences. Only the measure has.

This is not a quirk of incarceration data. The divergence between absolute and relative measures of disparity is a known formal result, demonstrated in health research where the two can move in opposite directions over the same period (Moonesinghe & Beckles, 2015), and the dependence of measured racial disparity on such choices has been shown directly for imprisonment (Sabol, Johnson, & Lynch, 2025). What the figure adds is a vivid sense of how cleanly the choice can reverse a headline: one decision, made before any modeling begins, has already settled whether democracy looks like a friend or an enemy of racial equality.

Demonstration 2: comparing states, or comparing a state with itself

Comparing different states with one another folds democracy together with everything else that distinguishes them, including region, demographics, history, and the local economy, and the first demonstration has already shown how badly that folding can distort an answer. A cleaner question

is longitudinal. When a single state's democracy rises or falls, does its incarceration move with it? Posed this way, each state serves as its own control, and the fixed differences between states drop out of the comparison.

To carry this out I split the democracy score into two parts. The first is each state's long-run average, which holds the between-state variation that the earlier figures relied on; the second is each year's deviation from that average, which holds the within-state variation that tracks change over time. The split follows Mundlak (1978). I then fit a linear mixed model with a random intercept for each state and a common year trend, using the lme4 package (Bates et al., 2015), so that the two kinds of variation are estimated side by side rather than confused with each other (Figure 3).

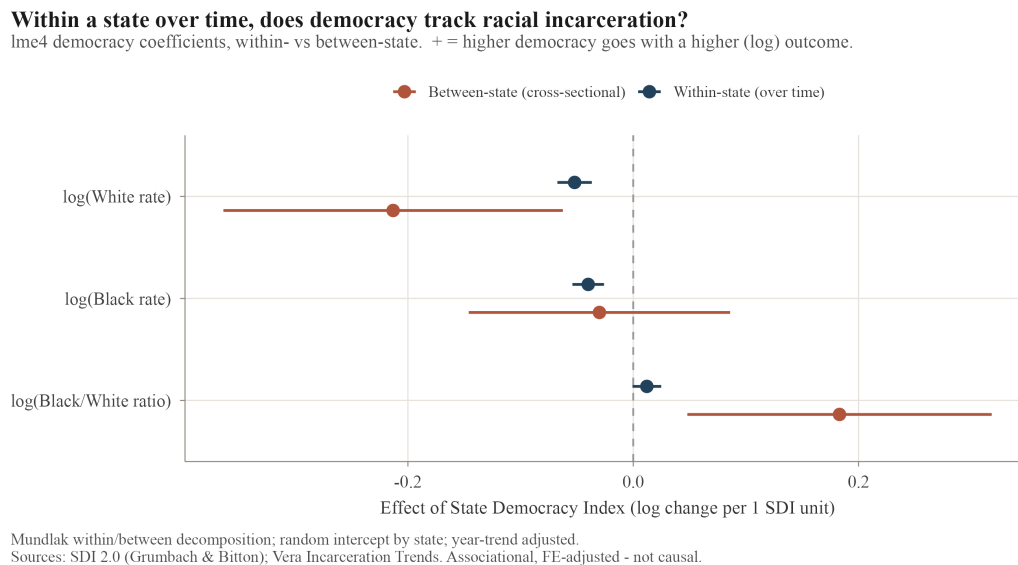


Figure 3: Within-state versus between-state democracy coefficients. The cross-sectional ratio effect (orange) disappears within states (blue).

The contrast is clean. The alarming cross-sectional result for the ratio turns out to be entirely between-state, at about +0.18, and to vanish within states, at about +0.01. In other words, it reflects which states are which, not what happens when a given state's democracy changes. Within states,

stronger democracy goes with lower incarceration of both groups, by roughly -0.04 for the Black rate and -0.05 for the white rate in log points per index point, and because the two rates fall together the ratio barely moves, exactly as the first demonstration would predict. The largest term in every model, however, is neither within nor between democracy but the year trend: the ratio fell about 2.6 percent a year and the Black rate about 2.7 percent, a nationwide decline that has little to do with state democracy. An analysis that left time out would mistake that tide for its own discovery. So the within-state link is real and consistent in direction, but it is modest, and the paper is stronger for saying so. One choice of comparison, cross-section against within-state, has already turned a striking effect into a small one.

Demonstration 3: a method that reads in more than the data hold

An average describes the typical state but conceals the differences between states, and those differences are part of the story. So the next question is whether states fall into recognizable types according to how democracy and incarceration moved together. Group-based trajectory modeling, developed in criminology by Nagin and Land (1993) to classify offending over the life course, is built for exactly this. Using the `gbmt` package (Magrini, 2022), it sorts the fifty states by their joint path from 2000 to 2022 on two series at once, democracy and the logged Black incarceration rate, and the BIC selects five well-separated classes, with an average assignment probability near 0.99 (Figure 4).

State trajectory classes (gbmt), 2000-2022: $ng = 5$

Class-mean paths. States typed jointly on democracy + log Black incarceration; ratio shown for the equity read.




Figure 4: Five state trajectory classes from the gbmt model. Each line is a class-mean path for democracy, Black incarceration, and the Black/White ratio.

The classes arrange themselves along a single spectrum. States that grew more democratic decarcerated the most, while those that backslid the furthest, the group Grumbach (2023) calls the laboratories of backsliding, reduced Black incarceration the least and even saw their white rate climb. This is the same relationship as before, now drawn as a contrast between recognizable groups of states. The method invites overreading, though, and resisting that is the lesson here. Every class decarcerated, because the national decline runs through all of them, so the groups differ in degree rather than in direction. The ratio converged toward four or five across the board. And the BIC keeps improving as classes are added, so the figure of five is not firmly pinned down; highly autocorrelated series like these can produce classes that are not really there. The classes are useful summaries, not natural kinds, and they remain entangled with region. A method can hand back a clean typology that looks more solid than the data warrant, and choosing to read it as description rather than discovery is itself part of the analysis.

Demonstration 4: a reform by its text, or by its effect

The state-level patterns are associations, not causes, and to get closer to cause I turn to a natural experiment. In 2018 Florida passed Amendment 4, restoring the vote to about 1.4 million people with felony convictions, the largest such restoration in modern U.S. history. With only one treated state the appropriate tool is synthetic control (Abadie, Diamond, & Hainmueller, 2010), implemented with the tidysynth package (Dunford, 2023). The method assembles a synthetic Florida from a weighted combination of states that did not change their felon-voting laws, chosen so that the combination reproduces Florida's turnout history before the reform, and then reads the reform's effect as the gap that opens afterward. The outcome is VAP turnout, whose denominator does not move when people are re-enfranchised, and a placebo test provides inference by asking how unusual Florida's gap is against the same gaps computed for every donor state.

Florida vs. synthetic Florida: VAP turnout, 2000-2022

Amendment 4 effective 2019 (first election 2020); SB7066 reimposed fines-first mid-2019




Figure 5: Florida compared with synthetic Florida on VAP turnout. The observed series (grey) tracks the synthetic series (pink) before and after the reform.

Placebo test: Florida's turnout gap vs. clean donor states




Figure 6: Placebo test. Florida's post-treatment gap (pink) sits inside the cloud of donor-state placebos.

The pre-reform fit is close, with a 2016 gap of only 0.07 points, which lends the comparison credibility (Figure 5). After the reform, the gaps are -1.4 points in 2020 and -4.8 in 2022, both negative rather than positive, and the placebo test returns a p-value of about 0.14, meaning Florida is not unusual among the donor states (Figure 6). Amendment 4 produced no measurable increase in turnout.

A null result is only worth reporting if its source is visible, and here it is. The reform barely took effect. Florida's disenfranchised share fell only from 1.25 percent of the voting-age population in 2018 to 1.16 percent in 2020 and 1.14 percent in 2022, a nearly flat line rather than the steep drop that re-enfranchising 1.4 million people would imply. The explanation is SB7066, which required people to pay all outstanding court fines and fees before regaining the vote, a requirement that works as a wealth test on the franchise and falls hardest on poor and disproportionately Black citizens. By its text the reform enfranchised more than a million people; by its effect it enfranchised few. Which

description counts as true depends on whether one reads the statute or the turnout, and here, once again, the framing is the finding.

Discussion

The four demonstrations make a single argument. In each, a routine decision settles the result: a measure of inequality reverses the sign, a comparison design shrinks a striking effect to a small one, a clustering method offers more structure than the data hold, and a reform reads as sweeping or negligible depending on whether one consults its text or its effect. The substantive conclusions matter, but they are secondary to this. On questions like these, measurement, design, and framing are not technicalities to be settled before the real work begins; they are where the answer is actually made.

That measurement and framing carry this weight is, as the background made clear, a well-worn point in the literature. What is easy to lose, stated only in the abstract, is how concretely and how quickly the choices bite, and how they accumulate. Here they are shown one after another, with numbers, on a case where the stakes are real: the most-cited measure of racial disparity in incarceration points the wrong way for a mechanical reason, and a celebrated voting-rights reform reads as a triumph or a non-event depending on the frame. For anyone who studies inequality, the practical counsel is cheap to follow. Carry more than one measure. Separate within from between. Treat data-driven groups as descriptions rather than kinds. Judge a policy by its effect rather than its text. And when the measures disagree, read the disagreement as a finding rather than a nuisance. Stated plainly, the substantive takeaways are modest. State democracy and racial inequality in

punishment are linked only weakly, and only within states over time; the link is small beside a national decline in incarceration; and re-enfranchisement on paper can be undone by a fee. None of these is a breakthrough. Taken together, though, they show how easily the convenient answer turns out to be wrong.

Several limits bound the claims. Aggregate turnout is a blunt instrument: even a fully effective Amendment 4, involving about 8 percent of the voting-age population, most of whom never register or vote, might move statewide turnout by less than a point, below what this design can detect, so the null means there was no large aggregate effect rather than that re-enfranchisement is pointless, and a sharper test would follow the formerly incarcerated individually in voter files. The first three demonstrations are observational rather than experimental, and they support the framing of the question rather than a causal answer to it. State democracy and incarceration are themselves estimated quantities. And the synthetic control rests on a single treated state, an imperfect midterm pre-fit, a small pool of clean donor states, and a 2020 treatment year disrupted by the pandemic.

Conclusion

This paper has used one case, race and democracy and punishment in the U.S. states, to make a methodological point: that the answer depends on how we measure and frame the question, often more than on the data. Change the measure and a disparity reverses; change the comparison and an effect shrinks; change the frame and a reform that freed a million voters did almost nothing. The lesson is not that the data are useless, but that the choices made around them deserve as much scrutiny as the findings themselves. On the questions of who is punished and who may vote, that scrutiny is not optional.

References

- Abadie, A., Diamond, A., & Hainmueller, J. (2010). Synthetic control methods for comparative case studies: Estimating the effect of California's tobacco control program. *Journal of the American Statistical Association*, 105(490), 493–505.
- Bates, D., Mächler, M., Bolker, B., & Walker, S. (2015). Fitting linear mixed-effects models using lme4. *Journal of Statistical Software*, 67(1), 1–48.
- Beck, A. J., & Blumstein, A. (2018). Racial disproportionality in U.S. state prisons: Accounting for the effects of racial and ethnic differences in criminal involvement, arrests, sentencing, and time served. *Journal of Quantitative Criminology*, 34(3), 853–883.
- Cesario, J., Johnson, D. J., & Terrill, W. (2019). Is there evidence of racial disparity in police use of deadly force? Analyses of officer-involved fatal shootings in 2015–2016. *Social Psychological and Personality Science*, 10(5), 586–595.
- D'Ignazio, C., & Klein, L. F. (2020). *Data feminism*. MIT Press.
- Dunford, E. (2023). *tidysynth: A tidy implementation of the synthetic control method* [R package].
- Gillborn, D., Warmington, P., & Demack, S. (2018). QuantCrit: Education, policy, “Big Data” and principles for a critical race theory of statistics. *Race Ethnicity and Education*, 21(2), 158–179.
- Grumbach, J. M. (2022). *Laboratories against democracy: How national parties transformed state politics*. Princeton University Press.
- Grumbach, J. M. (2023). Laboratories of democratic backsliding. *American Political Science*

Review, 117(3), 967–984.

Grumbach, J. M., & Bitton, F. (2024). *State Democracy Index 2.0* [Data set]. Democracy Policy Lab, University of California, Berkeley.

Kjellsson, G., Gerdtham, U.-G., & Petrie, D. (2015). Lies, damned lies, and health inequality measurements: Understanding the value judgments. *Epidemiology*, 26(5), 673–680.

Magrini, A. (2022). *gbmt: Group-based multivariate trajectory modeling* [R package].

Manza, J., & Uggen, C. (2006). *Locked out: Felon disenfranchisement and American democracy*. Oxford University Press.

McDonald, M. P. (2023). *United States Elections Project: Voter turnout data, 1980–2022* [Data set]. UF Election Lab, University of Florida.

Merry, S. E. (2016). *The seductions of quantification: Measuring human rights, gender violence, and sex trafficking*. University of Chicago Press.

Moonesinghe, R., & Beckles, G. L. A. (2015). Measuring health disparities: A comparison of absolute and relative disparities. *PeerJ*, 3, e1438.

Mundlak, Y. (1978). On the pooling of time series and cross section data. *Econometrica*, 46(1), 69–85.

Nagin, D. S., & Land, K. C. (1993). Age, criminal careers, and population heterogeneity: Specification and estimation of a nonparametric, mixed Poisson model. *Criminology*, 31(3), 327–362.

Quinn, D. M. (2025). Experimental effects of “opportunity gap” and “achievement gap” frames.

Sociology of Education. <https://doi.org/10.1177/00380407251321372>

Sabol, W. J., Johnson, T. L., & Lynch, J. P. (2025). Discrepancies in measures of racial and ethnic disparities: Implications for research, policy, and practice. *American Journal of Criminal Justice*. <https://doi.org/10.1007/s12103-025-09810-1>

The Sentencing Project. (2024). *Locked out 2024: Four million denied voting rights due to a felony conviction*.

van der Bles, A. M., van der Linden, S., Freeman, A. L. J., & Spiegelhalter, D. J. (2020). The effects of communicating uncertainty on public trust in facts and numbers. *Proceedings of the National Academy of Sciences*, 117(14), 7672–7683.

Vera Institute of Justice. (2024). *Incarceration trends* [Data set].